

I'm a hacker, and I love to build stuff for the Web.

TWITTER ★ LANYRD ★ GITHUB ★ G+ ★ EMAIL ★ MEAT

← ALL POSTS

## OPENPGP FOR COMPLETE BEGINNERS

Thursday 20<sup>th</sup> August, 2009

This is a rather comprehensive post detailing how to set up PGP encryption on your machine using GnuPG, an open-source implementation of the OpenPGP standard. I'll go through the core concepts of OpenPGP, generating a keypair and revocation certificate, backing up your configuration and data, and carrying out encryption and signing easily from both the command-line and your favourite-mail client.

Please note: the installation instructions are mostly geared towards Mac OS X users, but the command-line interaction and concepts are applicable to all systems.

### A QUICK INTRODUCTION TO ENCRYPTION

There are two main types of electronic encryption: *symmetric* and *asymmetric*. In symmetric encryption, a single, shared key is used to encrypt a document from cleartext to ciphertext, and the same key is used to decrypt from ciphertext back to cleartext. Obviously there is a problem using this to encrypt e-mails, since everyone has to know the shared key in order to read your e-mail, and if the key is public then it offers you zero protection (since any third party might be able to read the 'encrypted' data). Symmetric encryption is only worth anything if *only* the sender and receiver know the secret. It's also useful when there is *only* one party involved; for example, if you are encrypting files on your hard drive that only you will see, or if a website is encrypting users' data to put into a session cookie, etc.

process. The way this form of encryption works is that documents are encrypted with one key that can *only* be decrypted using the other key (they cannot even be decrypted with the same key that was used to encrypt them). So if I have keys A and B, which constitute a keypair, and I encrypt a document (such as an e-mail) with key A, it can *only* be decrypted by someone with key B, and vice versa. Algorithms which implement this kind of encryption are known as *transposition algorithms*, since a single key only works in one direction. In addition, another property of asymmetric encryption is that if you know for certain that key B decrypts a piece of ciphertext, key A *must* have been used to encrypt it.

### CONFIDENTIALITY

The way asymmetric encryption is used in maintaining confidentiality (i.e. the privacy of your precious documents) goes something like this: you generate a 'keypair' (a pair of keys such that documents encrypted with one can be decrypted with the other, and vice versa). One of these is your *public* key, the other is your *private* key. So, you share your public key with the world and keep your private key well, private. People who want to send you an encrypted e-mail find your public key, write the e-mail, encrypt it using the public key and send it to you. You, on receiving the e-mail, whip out your private key, decrypt the ciphertext and read it. This is how **confidentiality** is obtained with OpenPGP, and it's relatively simple.

In actual fact, it doesn't quite work like this. As it happens, and quite unfortunately so, asymmetric encryption has a few drawbacks. It is quite computationally intensive, and some algorithms can produce ciphertext up to twice the size of the original text. For these reasons, a technique known as *hybrid encryption* is employed. Rather than just encrypting a document with your public key, a unique, random *session key* is generated for that **number of documents**. This is also sometimes referred to as a **nonce** (needs to be completely random, essentially 'unguessable'; this involves having a good source of randomness (a.k.a. entropy) is very important in system security).

The document is encrypted with your public key. This will probably be much shorter than the document itself, so the computational load of asymmetric encryption is lightened significantly. The encrypted session key is attached to the symmetrically-encrypted document, and the session key is discarded forever. Upon receiving the encrypted document, you use your private key to decrypt the session key, and then use the session key to decrypt the entire document.

### AUTHENTICITY

One of the other uses of asymmetric encryption is for verifying the authenticity of data. For example, how do I know that this e-mail was *actually* sent to me from Joe Bloggs and not some impostor? We can take advantage of the fact that *if* our key can decrypt a piece of ciphertext, the other key must necessarily have encrypted it. First, the sender will take the e-mail and construct a *hash* using a *cryptographic hash function*. This function must have several properties:

- It should be impractical to find two distinct messages with the same hash.
- It should be impractical to obtain the original message given a hash (a CHFI is a *one-way function*).
- It should be easy to compute the hash value for any given message.

The sender then takes that hash and his/her own secret key, and *encrypts* the hash with the secret key. This encrypted hash, known as the signature, is bundled along with the message when it is sent. The recipient can then take the signature, the original message, and the sender's public key, and attempt to decrypt the signature using the public key. The recipient will hash the message themselves, and compare the decrypted signature to the hash of the message. If equal, the recipient can know for certain that the message was created by someone with a certain private key (without knowing the key themselves). This shows **authenticity** is obtained using OpenPGP.

It also helps to enforce data **integrity**, since the hash check will fail if the document has been changed in transit, so either data corruption or malicious intent will be detected.

This is usually used in tandem with encryption, providing confidentiality and integrity, but it's also pretty useful for things like software releases, where the maintainer will sign the release file with his/her private key. You'll download the release file and the corresponding `.sig` or `.asc` file, and verify that the package was downloaded intact and that it was published by the bona fide maintainer.

### KEYSERVERS

Obviously, if you have your public/private keypair, you'll want to somehow publish the public key somewhere where everyone can access it. Keyservers are a relatively simple way to do so. A keyserver is literally just a repository of public keys; I'll show you later how to

once). The session key

symmetrically with that session key, and then the session key is

push to and fetch from a keyserver,since this is an important way of keeping your keyring (i.e. the public keys of others you have saved locally) up-to-date, and also keeping others' keyrings up-to-date (such as when you make a change to your key).

WEB OF TRUST

One of the most important concepts in OpenPGP is that of a Web of Trust. Let's say Alice decides to generate a keypair using your identity (name and e-mail address). How are people to decide which belongs to who? Not only might someone send you e-mail encrypted for Alice, rendering you unable to read it, but Alice could also read that e-mail, constituting a major breach of security (not to mention a serious inconvenience). Alice could also sign documents, tricking recipients into thinking that you created them.

Fortunately, using the concept of signing as detailed above, you can sign a public key, as if to say "I trust this key—owner relationship." The more keys you sign that key, the more that key becomes "trusted". Eventually, you will have what is known as a "Web of Trust", where key A is signed by key B and C, key B is signed by keys A and D, and key C is signed by keys A, B, and D, for example. The set of keys and signatures makes up the Web of Trust. You can determine whether an identity is to be trusted or not by following the chain of trusted friends. This way, anyone trying to pose as a user will not be able to get much trust; you'll always go with the key that works out as most trusted.

I'll give an example. Let's say you and Bob are best friends. You trust him completely, so you've signed his key. Bob has a friend called Chris, who trusts completely, so he has signed his key. If Chris sends you a signed e-mail, you're very likely to believe that signature indeed belongs to Chris, because the authenticity of the key with which it is signed has itself been attested to by your best friend Bob (and you can verify that by checking Bob's signature against his public key). It can get confusing for even relatively small networks, but fortunately a simple computer algorithm can calculate a "trust level" for a given signature, based on following the sequence of connections of trust linking you to that signature.

You might also be happy to know that OpenPGP is a good reason to have a party. At [here](#).

OpenPGP key has a unique identifier: the morning after, everyone goes online and signs the keys they accepted the night before.

SUBKEYS

Subkeys are quite a subtle but powerful concept. The problem with the keypair is this: what happens when a key is compromised? You'll have to revoke your old keypair (a process which involves you verifying that you are the owner of the key, after which it will be marked as revoked), thus losing your entire Web of Trust, and then re-generate a key and publish that. Clearly not a very good situation to be in.

OpenPGP uses a concept known as a subkey to mitigate this problem. Rather than have a single keypair, you'll have a "master" keypair which acts as your identity, and a "subkey" pair which does the actual "cryption". If for some reason a subkey pair becomes compromised (i.e. someone finds out your passphrase), you can revoke just that subkey and create a new one, without losing your identity or Web of Trust.

GETTING STARTED WITH OPENPGP

This'll just be a brief guide to getting GnuPG installed and running on your Mac. I haven't included full instructions for GNU/Linux, BSD, other good OSes or Windows here because I'm not running any of those platforms right now; you can probably find some good guides to installing GnuPG elsewhere. Since at least this involves command-line work, which is likely to be the same on most platforms, it may be applicable to you anyway.

INSTALL GnuPG2

This is pretty easy to do on Mac: download the ZIP file from [here](#), unzip, install the .mpkg file, and you're done.

Alternatively, if you use MacPorts, you can just run `sudo port install gnuPG2`

Systems with `apt-get`, `rpm`, `dnf`, `zypper`, or `dnf` can also use their respective package managers to obtain a copy of GnuPG2.

GENERATE A KEYPAIR

This will involve using the command line. Launch Terminal.app (or your preferred terminal emulator) and do this:

```
gpg --gen-key
```

You'll see:

```
gpg (GnuPG/MacOSX 2.0.11); Copyright (C) 2009 Free Software Foundation, Inc.
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.

Please select what kind of key you want:
(1) RSA and RSA (default)
(2) RSA and ElGamal
(3) RSA (sign only)
(4) RSA (cert only)
Your selection? 1
```

Hit enter, since the defaults tend to work fine. Then:

```
RSA keys may be between 1024 and 4096 bits long.
What keysize do you want? (1024)
```


I personally go with 4096 at this point, but you do need to make a decision here. Whilst longer keys are more secure, they can also take a long time to generate (especially on older hardware; 4096-bit keys can take up to an hour to be generated on a PowerPC-based Mac), and "cryption" will also be slower. This might not matter if you and all your contacts are running relatively recent and powerful hardware, but once OpenPGP support for embedded and mobile devices becomes more available, it will present a serious issue. For this reason you may want to go for the default. After deciding, you'll be presented with this screen:

```
Requested keysize is 4096 bits
Please specify how long the key should be valid.
0 = key does not expire
inf = key expires in 0 days
0m = key expires in 0 minutes
0h = key expires in 0 hours
0d = key expires in 0 days
0w = key expires in 0 weeks
0Y = key expires in 0 years
key = key expires in 0 years
key is valid for? (0)
```

Hit enter again. This should be fine: if somehow your key becomes compromised you'll be able to revoke it using the revocation certificate (more on that later). It'll also ask you for confirmation; just hit `y` then enter again. The generator will then ask you to answer a series of questions:

```
Name needs to construct a user ID to identify your key.
Real name: (Type your full name, hit enter)
Email address: (Type your email address, hit enter)
Comment: (Optional) Type a comment to your homepage URL, hit enter
You selected 0 (0)
```

change their key fingerprints (since each





that and it'll get the correct key for you. At each stage the command will print all the recipients you've specified so far. To finish, just hit enter on a prompt without typing anything. Alternatively, you can specify all of the User IDs on the command line, like so:

```
gpg --encrypt document.txt --recipient 'A' --recipient 'B'
```

Where `--` is short for `--recipient`. You can pass `--` as many times as you like in one command.

PGP2 will encrypt the document and save the encrypted ciphertext as `document.txt.gpg` (mutatis mutandis). This file will, however, be in a binary format. If you need a plaintext version, just pass the `--armor` flag to the `pgp2` invocation. This produces an ASCII-armored version of the ciphertext called `document.txt.asc(m.m.)`, suitable for inclusion in e-mails, web pages etc.

To decrypt a document that someone has sent you, you can just run:

```
gpg --decrypt document.txt.gpg
```

This will prompt you for your private key's passphrase, and save the decrypted document in `document.txt(m.m.)`. It doesn't matter if the document is in an ASCII-armored or binary format, since GPG2 will recognize it either way.

SIGNING AND VERIFYING DOCUMENTS

Signing a document to prove its authenticity is also simple with GPG2. You don't need to specify any recipients, but you will need to provide your private key's passphrase. To sign a document, just run:

```
gpg --sign document.txt
```

You can also pass in `--armor`, if you need plaintext output. This will create a file called `document.txt.asc` depending on whether or not you armor the output. To verify a signature, just run:

```
gpg --verify document.txt.asc
```

For this to work, `document.txt` has to be in the same directory as its signature, and you also must have imported the public key of the person who signed it (since the public key is used in verification).

MAKING BACKUPS

At this point you'll want to make regular back-ups of your keyring and associated files. This is easy enough; simply back-up the contents of the `.gnupg` directory in your home folder, and also create a single back-up of the revocation certificate you generated in the last step. You'll actually want to back-up all the files *not* ending in either `pgp-keys` or the tilde character (~). The first denotes a *full* lock which is processed on the system called `gpg-agent` uses, and files ending in a tilde are back-up files (which obviously you won't need if you are actually making a back-up). You probably have your own back-up strategy, but I will tell you how best to make backups so that third parties can't easily access your backed-up data.

I use (and recommend) a piece of software called `TrueCrypt` to secure all of my backups. TrueCrypt allows you to create an encrypted drive in a file, which can then be mounted by the application, read from/written to, and then unmounted again. To most eyes it will look like it is made of completely random data. You can also encrypt whole partitions, allowing you to create an encrypted USB stick or external hard drive.

I'd recommend that you use TrueCrypt to create the drive on which you carry out all your back-ups. I won't go into the details of setting up an encrypted drive, since I was able to figure out what to do on my first try (it is a GUI application after all).

In addition, another nifty trick you can use is to move the external drive, and just make it into a symbolic link like so:

```
ln -s /dev/sda1 /home/yourname/.gnupg/.gpgkey
```

When the drive is mounted, GnuPG2 will work fine; when ejected, trying to use GnuPG2 will just fail with an error (which is what you want it to do anyway).

USING OPENPGP WITH AN E-MAIL CLIENT

OpenPGP's main area of application is in securing e-mails. Most important and confidential information is still sent via e-mail, so you need to be able to quickly "crypt" your e-mails and verify that they are from trusted sources. In this respect, e-mail client integration is key. Fortunately, most mail clients have mature and readily-available OpenPGP plugins.

If you want to be able to encrypt, decrypt, sign and verify e-mails from within Apple Mail 3.0 (the Leopard version) in a very simple way, you can use the freeware *fortissafe GPGMail* software. The installation instructions on its homepage work fine; just walk through the installer and restart Mail, and you should have all the good PGP functionality right there in Mail.

Thunderbird users can install the *Enigmail plugin*, to get PGP goodness from within. It can't offer much guidance, but I hear that it's pretty much the most complete solution available (and it works on all the platforms that Thunderbird does).

If you use Gmail from Firefox, you can install *FirePGP*, which will let you do all the useful OpenPGP operations from within your browser.

`.gnupg` directory to the encrypted